

Artistic Style Transfer

Suyan Qu

University of Wisconsin - Madison
Madison, WI 53706

squ27@wisc.edu

Abstract

Artistic style transfer is to apply the style of one image, the style image, typically a painting with special texture, to another image, the content image. The generated image will preserve the content presented in the content image while having the same style and texture of the style image. This project first demonstrates how to quantify the style and content difference between two images, through a pre-trained VGG19 model, which can be used to generate style transfered images through backpropagation, then presents a feed-forward deep network that can generate style-transfered images of size 300×400 in around 0.025 seconds.

1. Introduction

For long, the style and content of a image are considered as two distinct aspects of a picture. People associate the content of an image with objects captured by the image while the style of the image with how the captured object is presented, such as the color scheme, the stroke, etc. Despite that, the style and content of an image are actually two integrated parts that cannot be physically separated. However, modern technology makes this task achievable. With the help of deep learning neural networks, it is possible to apply the style of famous paintings like starry night to a present-day photo, faking Van Gogh's work in modern settings. This technology can potentially accelerate the advancement of contemporary art, as it allows people with no previous experience in painting to create visually appealing images. This project aims to implement such style transfer system using convolutional neural networks. First, a pre-trained image classification model will be used to quantify the style and content difference between two images [3], which can be used to generate style transfered images through backpropagation, then a feed-forward transformation model will be implemented based on the style and content difference defined [4].

2. Related Work

2.1. VGG 19

VGG 19 is a 19 layer convolutional neural network for large-scale image recognition [8] developed by Karen Simonyan and Andrew Zisserman in 2014. It can recognize images of 1000 different object categories. Its performance on ILSVRC-2012 dataset was remarkably well, with 23.7% top 1 validation error and 6.8% top 5 validation error.

2.2. Quantify Style and Content Difference

In 2015, Leon A. Gatys, Alexander S. Ecker, and Matthias Bethge proposed a method that quantifies the style and content difference between two images with the assistance of a pretrained neural network [3]. The basic intuition comes from the fact that shallower layers of an image classifier extract basic information of an input image, while deeper layers learn about the high level concepts that helps in classification. If two images have the same content, when they are sent into a image classifier, its deep layers should produce very similar results. Given two images, x and y and the output X^l and Y^l of a deep hidden layer l of a pre-trained network, content difference of these two images can be defined as

$$\mathcal{L}_{\text{content}}(x, y) = \frac{1}{2} \sum_{i,j} (X_{i,j}^l - Y_{i,j}^l)^2$$

On the other hand, the style of an image may have influence in all layers. A good measure of style should capture the texture information but not the global arrangement. One such measure is the gram matrix. Given input images x and y , and a hidden layer l of a pretrained image classifier, define gram matrices F^l and G^l as

$$F_{i,j}^l = \sum_k X_{i,k}^l X_{j,k}^l$$

$$G_{i,j}^l = \sum_k Y_{i,k}^l Y_{j,k}^l$$

Then for each layer of dimension $N \times M$, a style difference E_l can be calculated through

$$E_l = \frac{1}{4N_l^2 M_l^2} \sum_{i,j} (F_{i,j}^l - G_{i,j}^l)^2$$

With a chosen set of hidden layers L , the total style difference is the weighted sum of the style difference of each layer

$$\mathcal{L}_{\text{style}}(x, y) = \sum_{l \in L} w_l E_l$$

2.3. Feed-Forward Transformation Network

In recent years, various image transformation tasks have been trained with pixel-wise loss functions, such as semantic segmentation [6], depth and surface normal estimation [1], and etc. The feed-forward transformation model used in style transfer inspire from [6], where downsampling is used to reduce the dimension of feature map, then upsampling is applied to produce the final result.

2.4. Instance Normalization

In 2016, Ulyanov et al [10] point out that the result of stylization should not depend on the contrast of the content image in general as in batch normalization. In contrast, the style loss should be designed to transfer elements from a style image to the content image such that the contrast of the stylized image is similar to the contrast of the style image, and the general contrast of content images should be discarded. In another word, normalization should be specific to content images. In the recognition of this need, Ulyanov et al proposed instance normalization layer

$$y_{tijk} = \frac{x_{tijk} - \mu_i}{\sqrt{\sigma_i^2 + \epsilon}}$$

$$\mu = \frac{1}{HW} \sum_{l=1}^W \sum_{m=1}^H x_{tilm}$$

$$\sigma_{ti}^2 = \frac{1}{HW} \sum_{l=1}^W \sum_{m=1}^H (x_{tilm} - \mu_{ti})^2$$

3. Method

3.1. Backpropagation with Loss Function

Using the mechanism proposed by Gatys et al [3], given a style image s , a content image c , and a generated image g , a loss function can be calculated as a weighted sum of the style loss between style image and generated image, and the content loss between the generated image and the content image

$$\mathcal{L}_{\text{total}}(s, c, g) = \alpha \mathcal{L}_{\text{content}}(c, g) + \beta \mathcal{L}_{\text{style}}(s, g)$$



Figure 1. Sample outputs of the gradient descent approach. The upper left image is the original content image of a persian cat; the upper right is the original image styled with Hokusai; the lower left is the original image styled with Starry Night; the lower right is the original image styled with The Scream

With this loss function well defined, we can run an iterative algorithm similar to gradient descent on the generated image. During each iteration, the style, content, and generated images are sent into a pre-trained VGG 19 model. The *conv4_2* layer is used to calculate the content difference, which *conv1_1*, *conv2_1*, *conv3_1*, *conv4_1*, and *conv5_1* are used to calculate the style difference. With the loss function calculated, a gradient $\frac{\partial \mathcal{L}_{\text{total}}}{\partial g}$ can be calculated, and the generated image can be updated through gradient descent on the generated image

$$g' = g - \epsilon \text{sgn}\left(\frac{\partial \mathcal{L}_{\text{total}}}{\partial g}\right)$$

Since it is harder to recover the content details than the style, the generated image in the first iteration is the same as the content image, except some random noises are added to break the original style and accelerate the process. It is also worth mention that max pooling layers are replaced by average pooling layers for better results. Some sample output are presented in figure 1

3.2. Texture Synthesis

Gatys et al [2] also points out that the same algorithm can be used for texture synthesis. Texture synthesis, or texture extraction, is, as the name suggests, to generate a image with minor content but preserves the style of the original images. With the same algorithm above, a simple texture synthesis algorithm can be produced by running the backpropagation algorithm but replacing the content image with a blank image with noise. Some sample results are presented in figure 2.

3.3. Real-Time Style Transfer

The backpropagation algorithm, as pointed out by Johnson et al [4] and Ulyanov et al [9] [10], is extremely slow and inefficient. In practice, it takes, on a GTX 1050 ti,




Figure 2. Sample texture synthesis. The image on the left is the original image, *Starry Night*. On the right is the style synthesized from the texture synthesis algorithm




Figure 3. Sample output of the transform network. On the left is the input image, and on the right is the input image stylized with *Starry Night*

around 2 minutes to generate a 400×300 image, while generating a high resolution 1200×900 image can take 30 to 40 minutes. An alternative solution is to train a feed-forward transform network that takes in a content image and generates a stylized image. Such a model consists of two neural networks. A transform network f_W , that will transfer the content image to a stylized image, and a VGG19 model ϕ , used to calculate $\mathcal{L}_{\text{total}}$. Given the content image c and style image s , the generated image from this model would be $f_W(c)$. The loss function for training f_W would be the same as defined before:

$$W = \arg \min_W (\mathcal{L}_{\text{total}}(s, c, f_W(c)))$$

Notice that since the transform network takes in only the content image, the style image will be fixed during training, so the trained model will be specific to the style of that image used in training.

Our transform network is similar to the network implemented by Johnson et al [4] and Radford et al [7], except that the batch normalization layers are replaced by the instance normalization layers proposed by Ulyanov et al [10] to produce better results.

4. Experiments and Results

The model is trained on COCO 2014 train dataset [5] for 1000 epochs. Figure 3 shows a sample output from this model. As presented clearly from the sample output, the network does impressively well generating the style, but a lot of details of the content image are lost. One can only see the rough outline of the cat in the input image. In addition, when looking at the two windows to the left of the cat in the




Figure 4. On the left is the original output without warmup, and on the right is the output of the network trained with warmup. The same style and content image as in Figure 3 are used

input image, which appears to be very dark, the window on bottom is stylized with dark blue textures as expected, but the window on top is stylized with bright yellow textures. This unexpected and inconsistent shows that with random initialization, it is very hard to recover the content details compared with the texture. To add more bias to the content difference, in addition to increasing the weight of $\mathcal{L}_{\text{content}}$ in calculating $\mathcal{L}_{\text{total}}$, some warmup epochs are applied. During warmup, the loss function is defined to be the difference between the direct output from the neural network and the input content image. In this way, the transform network is expected to be an identity model and should generate the exact same image as input before starting training. Result of the network with warmup is shown in figure 4.

As clearly seen in this direct comparison, the result generated with warmup has more consistencies. The black segments in the background are stylized with dark textures, while the white carpet and the white fur of the cat in foreground are stylized with bright yellow textures. It also preserves more of the content details.

4.1. Quantitative Results

To evaluate the result of these models, the measure developed by Gatys et al [3] will be adopted again. For a content image c , a style image s , and a image generated from the model g , we will calculate $\mathcal{L}_{\text{content}}(g, c)$, which quantifies the content difference between content image and generated image, and $\mathcal{L}_{\text{style}}(g, s)$, which quantifies the style difference between style image and generated image. The smaller $\mathcal{L}_{\text{content}}(g, c)$ is, better the model preserves the content. The smaller $\mathcal{L}_{\text{style}}(g, s)$ is, the better the model stylizes the image. Table 1 shows these measures for the backpropagation model, the transform network without warmup, and the transform network with warmup. The table shows that the naive backpropagation method does impressively better than using transform network both in terms of style difference and content difference, and the use of warmup epochs effectively reduces the content difference by 8.4%, and reduces the style difference to only 11.6% of the original difference, which is kind of counter intuitive, as the warmup epochs are designed to improve

Model	$\mathcal{L}_{\text{content}}$	$\mathcal{L}_{\text{content}}$
Back-Prop	32143.725	2679063.8
Transform Net	46208.43	622746200.0
Transform Net + Warmup	42328.656	72501340.0

Table 1. Style and Content Difference using Different Approaches

Model	Runtime
Back-Prop 400 Iteration	27.03s
Transform Net	0.026s
Transform Net + Warmup	0.025s

Table 2. Runtime for 400×300 Images Using Different Model on Tesla K80




Figure 5. The output of transform network after warmup

the content difference.

Another Measure that worths notice is the runtime for these 2 methods, as shown in table 2. A transform network only need 0.025 seconds to transform a 400×300 image, which is 1081 times faster than the backpropagation method. Compared with the iterative gradient descent method, this method only requires a single pass of the input image into the network.

5. Conclusion

In this project, we implemented the gradient descent algorithm proposed by Gatys et al [3] and the feed-forward transform network proposed by Johnson et al [4] with some improvements including replacing the batch normalization layers by instance normalization [10] and warmup epochs. The results generative are impressive, but are not exactly what we expected. The output image from the transform network is of low resolution and many content details are lost. Figure 5 shows the output after the warmup epochs and one training epoch, implying that the content details may have lost during upsampling and downsampling in the first place.

Future explorations will be focusing on how to preserve more of the details of a content image in the transform network, including how to improve the resolution of the output. We will also try some different network architectures for this task to see their result.

References

- [1] David Eigen and Rob Fergus. Predicting depth, surface normals and semantic labels with a common multi-scale convolutional architecture. pages 2650–2658, 2015.
- [2] Leon Gatys, Alexander S Ecker, and Matthias Bethge. Texture synthesis using convolutional neural networks. pages 262–270, 2015.
- [3] Leon A Gatys, Alexander S Ecker, and Matthias Bethge. A neural algorithm of artistic style. *arXiv preprint arXiv:1508.06576*, 2015.
- [4] Justin Johnson, Alexandre Alahi, and Li Fei-Fei. Perceptual losses for real-time style transfer and super-resolution. pages 694–711, 2016.
- [5] Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C Lawrence Zitnick. Microsoft coco: Common objects in context. In *European conference on computer vision*, pages 740–755. Springer, 2014.
- [6] Jonathan Long, Evan Shelhamer, and Trevor Darrell. Fully convolutional networks for semantic segmentation. pages 3431–3440, 2015.
- [7] Alec Radford, Luke Metz, and Soumith Chintala. Unsupervised representation learning with deep convolutional generative adversarial networks. *arXiv preprint arXiv:1511.06434*, 2015.
- [8] Karen Simonyan and Andrew Zisserman. Very deep convolutional networks for large-scale image recognition. *arXiv preprint arXiv:1409.1556*, 2014.
- [9] Dmitry Ulyanov, Vadim Lebedev, Andrea Vedaldi, and Victor S Lempitsky. Texture networks: Feed-forward synthesis of textures and stylized images. 1(2):4, 2016.
- [10] Dmitry Ulyanov, Andrea Vedaldi, and Victor Lempitsky. Instance normalization: The missing ingredient for fast stylization. *arXiv preprint arXiv:1607.08022*, 2016.